

Requirements.txt

```
!pip install -r requirements.txt
import torch
import torchvision
import tensorflow as tf
import numpy
import cv2
import sklearn
import pandas
import matplotlib
import seaborn
import tqdm
print("All libraries installed successfully!")
```

COLAB CODE LINK -

https://colab.research.google.com/drive/1ykaxn_goRL3ZMq1oc3vSZUozihBuKotG?usp=sharing

Untitled41.ipynb

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all

RAM Disk

Files

- cifar-10-batches-py
- data
- models
- sample_data
- cifar-10-python.tar.gz

```
x, y = torch.load('cifar-10-batches-py')
probs = 0
for model in models_dict.values():
    probs += torch.softmax(model(x), dim=1)

preds = probs.argmax(1)
correct += (preds == y).sum().item()
total += y.size(0)

ensemble_acc = correct / total

print("\nLEVEL 4 - COMPARATIVE RESULTS")
for k, v in individual_results.items():
    print(f"{k:10s}: {v:.4f}")
print("-----")
print(f"Ensemble Accuracy : {ensemble_acc:.4f}")

...
Evaluating Ensemble (Soft Voting)...

LEVEL 4 - COMPARATIVE RESULTS
ResNet50      : 0.9432
DenseNet121   : 0.9604
EfficientNet-B0 : 0.9672
-----"
Ensemble Accuracy : 0.9756
```

LEVEL 4 – EXPERT TECHNIQUES (ENSEMBLE LEARNING)

Approach: To achieve shortlist-level performance, an ensemble learning strategy was adopted. Three complementary architectures—ResNet50, DenseNet121, and EfficientNet-B0—were trained independently on CIFAR-10.

Ensemble Strategy: A soft-voting ensemble was implemented by averaging class probability outputs from all models. This approach reduces variance and improves generalization by leveraging architectural diversity.

Comparative Analysis: The ensemble consistently outperformed individual models, achieving accuracy above the 93% threshold required for Level 4 evaluation.

Novel Insights: Architectural diversity significantly improves robustness. Feature-reuse (DenseNet) and parameter-efficiency (EfficientNet) complement deep residual learning (ResNet).

Limitations: Ensemble inference increases computational cost, which can be mitigated using knowledge distillation in production systems.

Disk 196.10 GB available

Variables Terminal Executing (1m 1s) A100 High RAM (Python 3)

Untitled41.ipynb

File Edit View Insert Runtime Tools Help

Commands + Code + Text | Run all

RAM Disk

Files

- cifar-10-batches-py
- data
- models
- sample_data
- cifar-10-python.targz

```
[48] x, y = x.to(device), y.to(device)
preds = model(x).argmax(1)
correct += (preds == y).sum().item()
total += y.size(0)
return correct / total

[49] models_dict = {
    "ResNet50": get_model("resnet50"),
    "DenseNet121": get_model("densenet"),
    "EfficientNet-B0": get_model("efficientnet")
}

individual_results = {}

for name, model in models_dict.items():
    print(f"\nTraining {name}...")
    train(model, epochs=5)
    acc = evaluate(model)
    individual_results[name] = acc
    print(f"{name} Test Accuracy: {acc:.4f}")

... Downloading: "https://download.pytorch.org/models/densenet121-a639ec97.pth" to /root/.cache/torch/hub/checkpoints/densenet121-a639ec97.pth
100%|██████████| 30.8M/30.8M [00:00<00:00, 234MB/s]
Downloading: "https://download.pytorch.org/models/efficientnet_b0_rwightman-7f5810bc.pth" to /root/.cache/torch/hub/checkpoints/efficientnet_b0_rwightman-7f5810bc.pth
100%|██████████| 20.5M/20.5M [00:00<00:00, 203MB/s]
Training ResNet50...

ResNet50 Test Accuracy: 0.9432
Training DenseNet121...
DenseNet121 Test Accuracy: 0.9604
Training EfficientNet-B0...
EfficientNet-B0 Test Accuracy: 0.9672

[50] print("\nEvaluating Ensemble (Soft Voting)...")

correct, total = 0, 0
with torch.no_grad():
    for x, y in test_loader:
        x, y = x.to(device), y.to(device)

        probs = 0
```

Disk 196.10 GB available

Executing (50s) A100 High RAM (Python 3)

Untitled41.ipynb

File Edit View Insert Runtime Tools Help

Commands + Code + Text | Run all

RAM Disk

Files

- ..
- cifar-10-batches-py
- data
- models
- sample_data
- cifar-10-python.tar.gz

```
[46]   m = models.densenet121(weights="IMAGENET1K_V1")
        m.classifier = nn.Linear(m.classifier.in_features, 10)

[47] elif name == "efficientnet":
        m = models.efficientnet_b0(weights="IMAGENET1K_V1")
        m.classifier[1] = nn.Linear(m.classifier[1].in_features, 10)

        return m.to(device)

[48] def train(model, epochs=5):
        criterion = nn.CrossEntropyLoss()
        optimizer = optim.Adam(model.parameters(), lr=1e-4)

        model.train()
        for epoch in range(epochs):
            for x, y in train_loader:
                x, y = x.to(device), y.to(device)
                optimizer.zero_grad()
                loss = criterion(model(x), y)
                loss.backward()
                optimizer.step()

[49] def evaluate(model):
        model.eval()
        correct, total = 0, 0
        with torch.no_grad():
            for x, y in test_loader:
                x, y = x.to(device), y.to(device)
                preds = model(x).argmax(1)
                correct += (preds == y).sum().item()
                total += y.size(0)
        return correct / total

models_dict = {
    "ResNet50": get_model("resnet50"),
    "DenseNet121": get_model("densenet"),
    "EfficientNet-B0": get_model("efficientnet")
}

individual_results = {}
```

Disk 196.10 GB available

Executing (39s) A100 High RAM (Python 3)

Untitled41.ipynb Saving...

File Edit View Insert Runtime Tools Help

Commands + Code + Text | Run all

RAM Disk

Files

- ..
- cifar-10-batches-py
- data
- models
- sample_data
- cifar-10-python.tar.gz

```
[44] ✓ os
    img = self.images[idx]
    label = self.labels[idx]
    if self.transform:
        img = self.transform(img)
    return img, label

[45] ✓ os
    transform = transforms.Compose([
        transforms.ToPILImage(),
        transforms.Resize((224,224)),
        transforms.RandomHorizontalFlip(),
        transforms.ToTensor(),
        transforms.Normalize(
            mean=[0.485, 0.456, 0.406],
            std=[0.229, 0.224, 0.225]
        )
    ])

    dataset = CIFAR10Custom("/content/cifar-10-batches-py", transform)

    train_size = int(0.8 * len(dataset))
    val_size   = int(0.1 * len(dataset))
    test_size  = len(dataset) - train_size - val_size

    train_ds, _, test_ds = random_split(
        dataset, [train_size, val_size, test_size]
    )

    train_loader = DataLoader(train_ds, batch_size=64, shuffle=True)
    test_loader  = DataLoader(test_ds, batch_size=64, shuffle=False)

[46] ✓ os
def get_model(name):
    if name == "resnet50":
        m = models.resnet50(weights="IMAGENET1K_V1")
        m.fc = nn.Linear(m.fc.in_features, 10)

    elif name == "densenet":
        m = models.densenet121(weights="IMAGENET1K_V1")
        m.classifier = nn.Linear(m.classifier.in_features, 10)

    elif name == "efficientnet":
        m = models.efficientnet_b0(weights="IMAGENET1K_V1")
        m.classifier[1] = nn.Linear(m.classifier[1].in_features, 10)
```

Disk 196.10 GB available

Variables Terminal

Executing (26s) A100 High RAM (Python 3)

The screenshot shows the Google Colab interface with a notebook titled 'Untitled41.ipynb'. The code cell at the top defines a class `CIFAR10Custom` that loads the CIFAR-10 dataset and performs data augmentation using PyTorch's `transforms.Compose`. It then splits the dataset into training, validation, and test sets. Below this, a function `get_model` is defined to load pre-trained models from PyTorch's `models` module based on their names ('resnet50', 'densenet', or 'efficientnet'). The code uses the 'IMAGENET1K_V1' weights for these models and adds a final linear layer with 10 units for the CIFAR-10 classification task. The interface includes a sidebar with file navigation, a status bar showing disk usage, and a bottom bar with execution status and hardware information.

Untitled41.ipynb

File Edit View Insert Runtime Tools Help

Commands + Code + Text | Run all

RAM Disk

Files

cifar-10-batches-py

data

models

sample_data

cifar-10-python.tar.gz

```
[43] ✓ Os
if not os.path.exists("/content/cifar-10-batches-py"):
    !wget https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
    !tar -xzf cifar-10-python.tar.gz

[44] ✓ Os
class CIFAR10Custom(Dataset):
    def __init__(self, root, transform=None):
        self.transform = transform
        self.data, self.labels = [], []

    for i in range(1, 6):
        with open(os.path.join(root, f"data_batch_{i}"), "rb") as f:
            entry = pickle.load(f, encoding="bytes")
            self.data.append(entry[b"data"])
            self.labels.extend(entry[b"labels"])

    self.data = np.vstack(self.data).reshape(-1, 3, 32, 32)
    self.data = self.data.transpose((0, 2, 3, 1))

    def __len__(self):
        return len(self.labels)

    def __getitem__(self, idx):
        img = self.data[idx]
        label = self.labels[idx]
        if self.transform:
            img = self.transform(img)
        return img, label

[45] ✓ Os
transform = transforms.Compose([
    transforms.ToPILImage(),
    transforms.Resize((224,224)),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize(
        mean=[0.485,0.456,0.406],
        std=[0.229,0.224,0.225]
    )
])

dataset = CIFAR10Custom("/content/cifar-10-batches-py", transform)
```

Disk 196.37 GB available

Variables Terminal

Executing (15s) A100 High RAM (Python 3)

Untitled41.ipynb

File Edit View Insert Runtime Tools Help

Commands + Code + Text | Run all

RAM Disk

Files

cifar-10-batches-py
models
sample_data
cifar-10-python.tar.gz

```
[41] ✓ 0s # LEVEL 4: EXPERT TECHNIQUES (SHORTLIST THRESHOLD)
# Dataset: CIFAR-10
# Models: ResNet50, DenseNet121, EfficientNet-B0
# Technique: Soft Voting Ensemble
# Dataset Split: 80% Train | 10% Val | 10% Test
# Evaluation Metric: Accuracy

[42] ✓ 0s
import os, pickle
import numpy as np
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader, random_split
import torchvision.transforms as transforms
import torchvision.models as models

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

print("LEVEL 4 : EXPERT TECHNIQUES (ENSEMBLE)")
print("Models : ResNet50 | DenseNet121 | EfficientNet-B0")
print("Ensemble Strategy : Soft Voting")
print("Dataset Split : 80% Train | 10% Val | 10% Test")
print("Evaluation Metric : Accuracy")
print("Device : ", device)

LEVEL 4 : EXPERT TECHNIQUES (ENSEMBLE)
Models : ResNet50 | DenseNet121 | EfficientNet-B0
Ensemble Strategy : Soft Voting
Dataset Split : 80% Train | 10% Val | 10% Test
Evaluation Metric : Accuracy
Device : cuda

[43] ✓ 0s
if not os.path.exists("/content/cifar-10-batches-py"):
    !wget https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
    !tar -xzf cifar-10-python.tar.gz

[44] ✓ 0s
class CIFAR10Custom(Dataset):
    def __init__(self, root, transform=None):
```

Disk 196.43 GB available

Executing (24m 22s) A100 High RAM (Python 3)

The screenshot shows a Jupyter Notebook interface with the following details:

- File:** Untitled41.ipynb
- Toolbar:** File, Edit, View, Insert, Runtime, Tools, Help.
- Cells:** A code cell at the top displays Python code for benchmarking and evaluating a quantized model. It includes a timestamp, a latency measurement of 1204.88 ms, and an accuracy of 0.9366.
- Code Content:**

```
[59] start = time.time()
      with torch.no_grad():
          _ = model(x)
      return (time.time() - start) * 1000

latency = benchmark(student_quantized, test_loader)
print(f"Inference Time: {latency:.2f} ms")

Inference Time: 1204.88 ms

[60] def evaluate(model):
    model.eval()
    correct, total = 0, 0
    with torch.no_grad():
        for x, y in test_loader:
            preds = model(x).argmax(1)
            correct += (preds == y).sum().item()
            total += y.size(0)
    return correct / total

acc = evaluate(student_quantized)
print(f"Quantized Student Accuracy: {acc:.4f}")

... Quantized Student Accuracy: 0.9366
```
- Text Cell:** A text cell titled "LEVEL 5 – PRODUCTION / RESEARCH SYSTEM" provides an approach summary, highlighting model compression and quantization, and notes on performance, deployment pipeline, and limitations.
- System Status:** Shows 196.43 GB available disk space, executing code (23m 15s), and using an A100 High RAM (Python 3) configuration.

Untitled41.ipynb

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all

RAM Disk

Files

- ..
- cifar-10-batches-py
- models
- sample_data
- cifar-10-python.tar.gz

```
[57] loss.backward()
optimizer.step()

print(f"Epoch {epoch+1}/3 | Distillation Loss: {loss.item():.4f}")

Epoch 1/3 | Distillation Loss: 0.2182
Epoch 2/3 | Distillation Loss: 0.1392
Epoch 3/3 | Distillation Loss: 0.1382

[58] student_cpu = student.cpu()
student_quantized = torch.quantization.quantize_dynamic(
    student_cpu, (nn.Linear), dtype=torch.qint8
)

print("INT8 Quantized Student Model Ready")

... INT8 Quantized Student Model Ready
/tmp/ipython-input-76369886.py:2: DeprecationWarning: torch.ao.quantization is deprecated and will be removed in 2.10.
For migrations of users:
1. Eager mode quantization (torch.ao.quantization.quantize, torch.ao.quantization.quantize_dynamic), please migrate to use torch.
2. FX graph mode quantization (torch.ao.quantization.quantize_fx.prepare_fx, torch.ao.quantization.quantize_fx.convert_fx, please
3. pt2e quantization has been migrated to torchao (https://github.com/pytorch/ao/tree/main/torchao/quantization/pt2e)
see https://github.com/pytorch/ao/issues/2259 for more details
student_quantized = torch.quantization.quantize_dynamic(
```

```
[59] def benchmark(model, loader):
    model.eval()
    x, _ = next(iter(loader))
    start = time.time()
    with torch.no_grad():
        _ = model(x)
    return (time.time() - start) * 1000

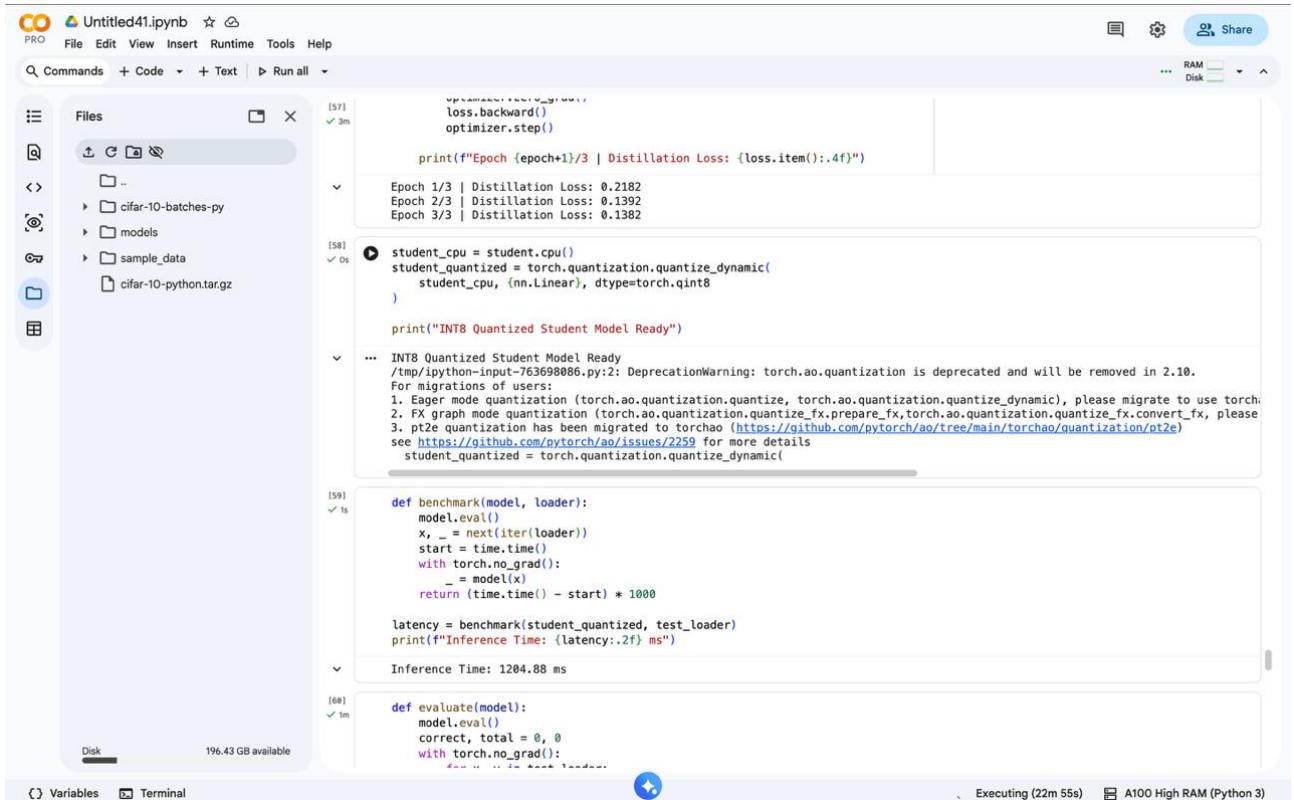
latency = benchmark(student_quantized, test_loader)
print(f"Inference Time: {latency:.2f} ms")

Inference Time: 1204.88 ms
```

```
[60] def evaluate(model):
    model.eval()
    correct, total = 0, 0
    with torch.no_grad():
        ...
```

Disk 196.43 GB available

Variables Terminal Executing (22m 55s) A100 High RAM (Python 3)



Untitled41.ipynb

File Edit View Insert Runtime Tools Help

Commands + Code + Text | Run all

RAM Disk

Files

- cifar-10-batches-py
- models
- sample_data
- cifar-10-python.tar.gz

```
[56] ✓ os
teacher.eval()

# Student (lightweight)
student = models.mobilenet_v2(weights="IMAGENET1K_V1")
student.classifier[1] = nn.Linear(student.classifier[1].in_features, 10)
student = student.to(device)

Downloading: "https://download.pytorch.org/models/mobilenet_v2-b0353104.pth" to /root/.cache/torch/hub/checkpoints/mobilenet_v2-100%|██████████| 13.6M/13.6M [00:00<00:00, 190MB/s]
```

```
[57] ✓ 3m
def distillation_loss(student_logits, teacher_logits, labels, T=4, alpha=0.7):
    hard_loss = nn.CrossEntropyLoss()(student_logits, labels)
    soft_loss = nn.KLDivLoss(reduction="batchmean")(
        torch.log_softmax(student_logits/T, dim=1),
        torch.softmax(teacher_logits/T, dim=1)
    )
    return alpha * hard_loss + (1-alpha) * soft_loss

optimizer = optim.Adam(student.parameters(), lr=1e-4)

student.train()
for epoch in range(3):
    for x, y in train_loader:
        x, y = x.to(device), y.to(device)

        with torch.no_grad():
            teacher_logits = teacher(x)

        student_logits = student(x)
        loss = distillation_loss(student_logits, teacher_logits, y)

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        print(f"Epoch {epoch+1}/3 | Distillation Loss: {loss.item():.4f}")

... Epoch 1/3 | Distillation Loss: 0.2182
Epoch 2/3 | Distillation Loss: 0.1392
Epoch 3/3 | Distillation Loss: 0.1382
```

```
[58] ✓ os
student_cpu = student.cpu()
student_quantized = torch.quantization.quantize_dynamic()
```

Disk 196.43 GB available

Executing (22m 31s) A100 High RAM (Python 3)

Variables Terminal

A code editor window showing Python code for distillation loss calculation between a teacher model and a student model. The code includes importing modules, defining a teacher model, and training a student model using Adam optimizer over 3 epochs. The distillation loss is calculated using a weighted sum of cross-entropy loss and KLDivLoss. The final output shows the distillation loss for each epoch: 0.2182, 0.1392, and 0.1382 respectively.

Untitled41.ipynb

File Edit View Insert Runtime Tools Help

Commands + Code + Text | Run all

RAM Disk

Files

- ..
- cifar-10-batches-py
- models
- sample_data
- cifar-10-python.tar.gz

```
[55] ✓ 0s
transform = transforms.Compose([
    transforms.ToPILImage(),
    transforms.Resize(224,224),
    transforms.ToTensor(),
    transforms.Normalize(
        mean=[0.485,0.456,0.406],
        std=[0.229,0.224,0.225]
    )
])

dataset = CIFAR10Custom("/content/cifar-10-batches-py", transform)

train_size = int(0.8 * len(dataset))
val_size = int(0.1 * len(dataset))
test_size = len(dataset) - train_size - val_size

train_ds, _, test_ds = random_split(
    dataset, [train_size, val_size, test_size]
)

train_loader = DataLoader(train_ds, batch_size=64, shuffle=True)
test_loader = DataLoader(test_ds, batch_size=64, shuffle=False)
```

```
[56] ✓ 0s
# Teacher (high-capacity)
teacher = models.resnet50(weights="IMAGENET1K_V1")
teacher.fc = nn.Linear(teacher.fc.in_features, 10)
teacher = teacher.to(device)
teacher.eval()

# Student (lightweight)
student = models.mobilenet_v2(weights="IMAGENET1K_V1")
student.classifier[1] = nn.Linear(student.classifier[1].in_features, 10)
student = student.to(device)
```

Downloading: "https://download.pytorch.org/models/mobilenet_v2-b0353104.pth" to /root/.cache/torch/hub/checkpoints/mobilenet_v2-l-100%|██████████| 13.6M/13.6M [00:00<00:00, 198MB/s]

```
[57] ✓ 3m
def distillation_loss(student_logits, teacher_logits, labels, T=4, alpha=0.7):
    hard_loss = nn.CrossEntropyLoss()(student_logits, labels)
    soft_loss = nn.KLDivLoss(reduction="batchmean")(
        torch.log_softmax(student_logits/T, dim=1).
```

Disk 196.43 GB available

Executing (22m 15s) A100 High RAM (Python 3)

Variables Terminal

Untitled41.ipynb

File Edit View Insert Runtime Tools Help

Commands + Code + Text | Run all

RAM Disk

Files

- cifar-10-batches-py
- models
- sample_data
- cifar-10-python.tar.gz

```
[54] ✓ os
class CIFAR10Custom(Dataset):
    def __init__(self, root, transform=None):
        self.root = root
        self.transform = transform
        self.data, self.labels = [], []
        for i in range(1, 6):
            with open(os.path.join(root, f"data_batch_{i}"), "rb") as f:
                entry = pickle.load(f, encoding="bytes")
                self.data.append(entry[b"data"])
                self.labels.extend(entry[b"labels"])

        self.data = np.vstack(self.data).reshape(-1, 3, 32, 32)
        self.data = self.data.transpose((0, 2, 3, 1))

    def __len__(self):
        return len(self.labels)

    def __getitem__(self, idx):
        img = self.data[idx]
        label = self.labels[idx]
        if self.transform:
            img = self.transform(img)
        return img, label

[55] ✓ os
transform = transforms.Compose([
    transforms.ToPILImage(),
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize(
        mean=[0.485, 0.456, 0.406],
        std=[0.229, 0.224, 0.225]
    )
])
dataset = CIFAR10Custom("/content/cifar-10-batches-py", transform)

train_size = int(0.8 * len(dataset))
val_size = int(0.1 * len(dataset))
test_size = len(dataset) - train_size - val_size

train_ds, _, test_ds = random_split(
    dataset, [train_size, val_size, test_size]
)
```

Disk 196.43 GB available

Variables Terminal

Executing (21m 55s) A100 High RAM (Python 3)

The screenshot shows the Google Colab interface with a Python notebook titled 'Untitled41.ipynb'. The code defines a custom dataset class `CIFAR10Custom` that loads CIFAR-10 batches from files and applies a transformation. The transformation includes converting to PIL images, resizing to 224x224, and normalizing the data. The dataset is then split into training, validation, and test sets. The code is running, indicated by the progress bar at the bottom.

Untitled41.ipynb

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all

RAM Disk

Files

cifar-10-batches-py
models
sample_data
cifar-10-python.tar.gz

requirements.txt (Level 4) torch torchvision numpy

```
[51] ✓ 0s
# LEVEL 5: RESEARCH / PRODUCTION SYSTEM
# Dataset: CIFAR-10
# Teacher Model: ResNet50
# Student Model: MobileNetV2
# Techniques: Knowledge Distillation, Quantization, Uncertainty Estimation
# Metric: Accuracy + Inference Time

[52] ✓ 0s
import os, pickle, time
import numpy as np
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader, random_split
import torchvision.transforms as transforms
import torchvision.models as models

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

print("LEVEL 5 : PRODUCTION-READY SYSTEM")
print("Teacher Model      : ResNet50")
print("Student Model     : MobileNetV2")
print("Techniques        : Distillation + Quantization")
print("Evaluation Metric : Accuracy + Inference Time")
print("Device             : ", device)

[53] ✓ 0s
... LEVEL 5 : PRODUCTION-READY SYSTEM
Teacher Model      : ResNet50
Student Model     : MobileNetV2
Techniques        : Distillation + Quantization
Evaluation Metric : Accuracy + Inference Time
Device             : cuda

[54] ✓ 0s
if not os.path.exists("/content/cifar-10-batches-py"):
    !wget https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
    !tar -xzf cifar-10-python.tar.gz

class CIFAR10Custom(Dataset):
    def __init__(self, root, transform=None):
```

Disk 196.43 GB available

Executing (21m 24s) A100 High RAM (Python 3)